

Variational autoencoders (VAEs) are generative models used in machine learning (ML) to generate new data in the form of variations of the input data they're trained on. In addition to this, they also perform tasks common to other autoencoders, such as denoising.